




Containers for Beginners

Michael Irwin - @mikesir87

Virginia Tech; Docker Captain




**Disclaimer: I cannot explain
sprinkle pods either!**

Quick History of Shipping



Source: <https://www.publicdomainpictures.net/en/view-image.php?image=275355>



Source: https://en.wikipedia.org/wiki/Rail_freight_in_Great_Britain



Source: <https://pxhere.com/en/photo/553345>




Software = Shipping?

Shipping in Software



Source: <https://www.usafe.af.mil/News/Photos/igphoto/2000887438/>



**Either of these two scenarios
sound familiar to you?**

Welcome!
Glad to have
you on the
team!

Clone the repo,
use the wiki for
setup instructions,
and update the
docs as needed.
Good luck!





Imagine if...

Creating Images

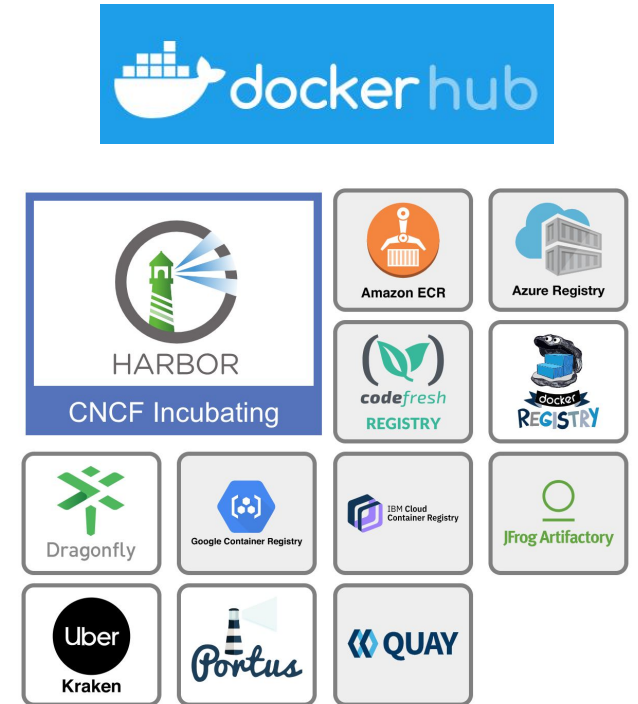
- Best practice is to use a [Dockerfile](#)
 - A text file that serves as a script to build an image
- Build using the [docker build](#) command

```
FROM node
WORKDIR /app
COPY package.json yarn.lock .
RUN yarn install
COPY src ./src
CMD ["node", "src/index.js"]
```




Sharing Images

- Once built, the image is only available locally
- To share, push it to a registry using `docker push`
 - Docker Hub is the default registry
 - Docker EE includes the Docker Trusted Registry
 - Many other third-party offerings available too
- Once shared, others can pull the image



Source: <https://landscape.cncf.io>



Let's build an image!

What's a container then?

- While a container looks like a VM, it isn't!
 - A container is **just another process** on the machine
- It uses namespaces and control groups (cgroups) to provide isolation
 - Namespaces include network, process, user, IPC, mount, and others
- To run a container, use the `docker container run` command



Containers vs VMs

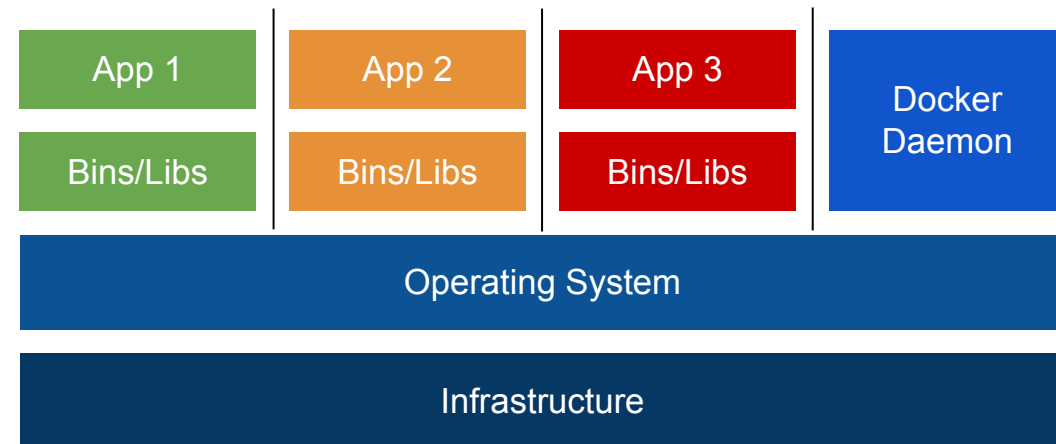
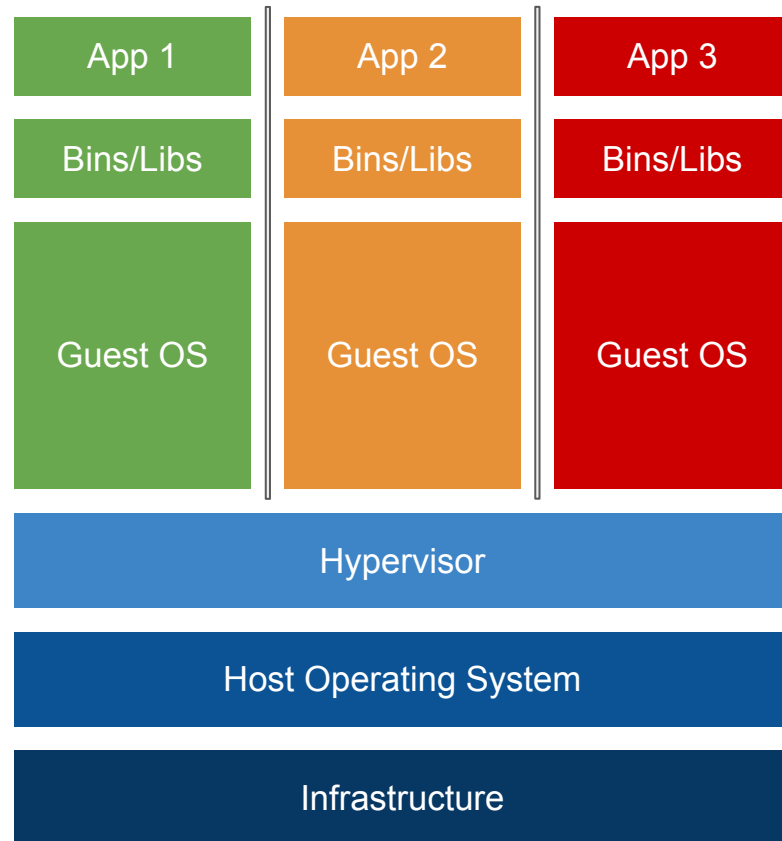


Image Layering

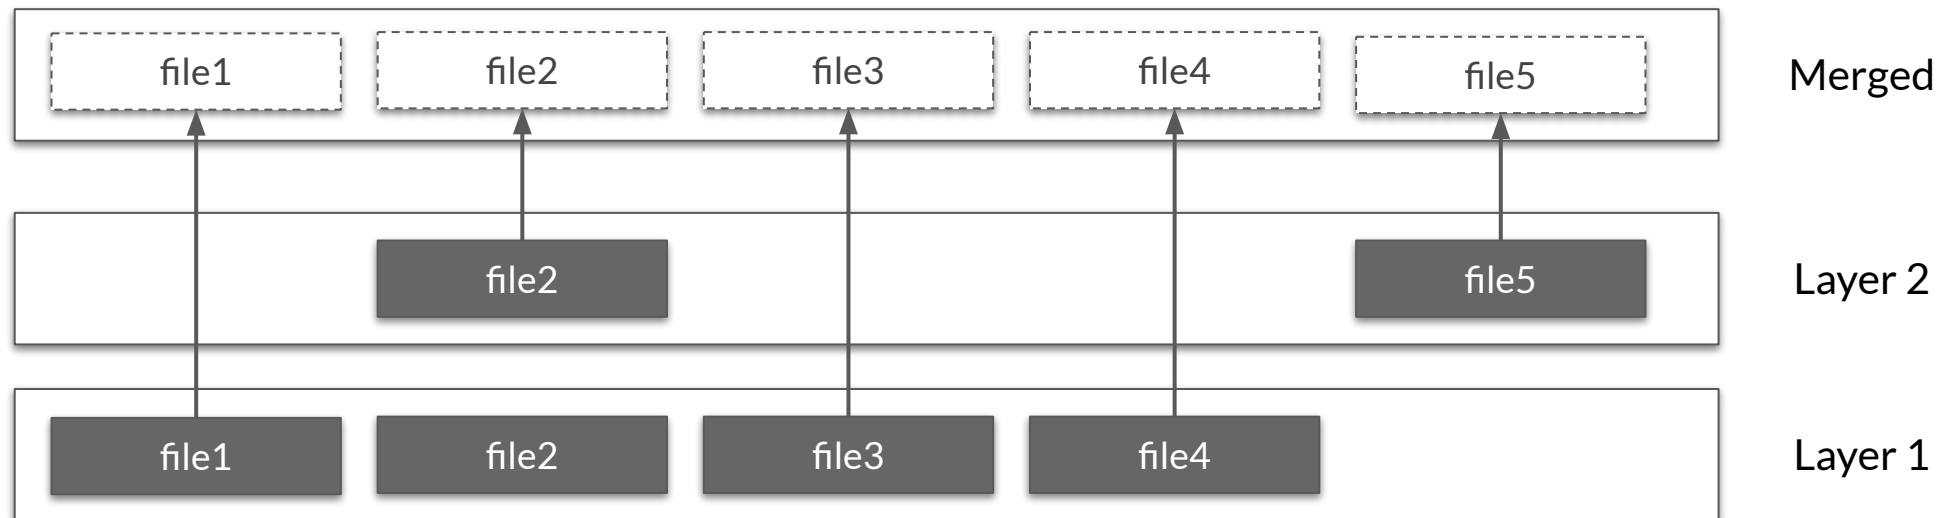
- Images are composed of layers of filesystem changes
 - Each layer can add or remove from the previous layer
 - Each layer's filesystem changes are stored as a single tar file
- Each command in a [Dockerfile](#) creates a new layer
- Use the [docker image history](#) command to see the layers and the command that was used to create each layer

```
~/dockercon docker image history mikesir87/mystery-image
```

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
4411b0d30bb7	9 months ago	/bin/sh -c #(nop) CMD ["/bin/sh" "-c" "npm ...	0B	
<missing>	9 months ago	/bin/sh -c #(nop) EXPOSE 3000	0B	
<missing>	9 months ago	/bin/sh -c npm install && rm /app/src/settin...	2.99MB	
<missing>	9 months ago	/bin/sh -c #(nop) COPY dir:ca65ca169aa7e7485...	656B	
<missing>	9 months ago	/bin/sh -c #(nop) WORKDIR /app	0B	
<missing>	9 months ago	/bin/sh -c #(nop) CMD ["node"]	0B	
<missing>	9 months ago	/bin/sh -c apk add --no-cache --virtual .bui...	4.51MB	
<missing>	9 months ago	/bin/sh -c #(nop) ENV YARN_VERSION=1.7.0	0B	
<missing>	9 months ago	/bin/sh -c addgroup -g 1000 node && addu...	61.9MB	
<missing>	9 months ago	/bin/sh -c #(nop) ENV NODE_VERSION=10.6.0	0B	
<missing>	9 months ago	/bin/sh -c #(nop) CMD ["/bin/sh"]	0B	
<missing>	9 months ago	/bin/sh -c #(nop) ADD file:6ee19b92d5cb1bf14...	4.2MB	


Layer contents

- Layers are unioned together to make a full filesystem
 - Each layer can add files as needed
 - Files in “higher” layers replace the same file in “lower” layers
- The container uses the “merged” view



What about deleted files?

- Deleted files are represented in a layer as a “whiteout” file
- Whiteout files are only used by the filesystem driver and not visible in the merged filesystem



WARNING!

**Be careful what you put into images.
Deleted files might not actually be
gone!**

Two Best Practices Incoming!

Clean up as you go!

- Don't wait until the end of the Dockerfile to “clean” up
- Chain **RUN** commands together to clean things as you go

```
FROM ubuntu
RUN apt-get update
RUN apt-get install -y python python-pip
RUN pip install awscli
RUN apt-get autoremove --purge -y python-pip
```

**Net change of image size from
512MB to 183MB (64% reduction)**

```
FROM ubuntu
RUN apt-get update && \
  apt-get install -y python python-pip && \
  pip install awscli && \
  apt-get autoremove --purge -y python-pip && \
  rm -rf /var/lib/apt/lists/*
```

Keep images tight and focused

- Only install the deps/tools/packages that are necessary
- Use multi-stage builds to separate build-time and run-time dependencies

```
FROM node AS build
WORKDIR /usr/src/app
COPY package.json yarn.lock .
RUN yarn install
COPY public ./public
COPY src ./src
RUN yarn build

FROM nginx:alpine
COPY nginx.conf /etc/nginx/nginx.conf
COPY --from=build /usr/src/app/build /usr/share/nginx/html
```

Sample multi-stage build for a React app

How do you persist data?

Volumes

- Volumes provide the ability to persist/supply data
- Bind mount volumes
 - You choose where to persist the data
 - Example: `-v $HOME/mysql-data:/var/lib/mysql`
- Named volumes
 - Let Docker choose where to persist the data
 - Can use `docker volume inspect` to find actual location
 - Example: `-v mysql-data:/var/lib/mysql`

Show me these volumes!




Docker Compose

- Makes defining and running multi-container apps super easy
- Uses a YAML file for configuration ([docker-compose.yml](#))
 - Often included in project source repo at the root of the project
- With a single command, start all containers/services for an app
- Tool is bundled with Docker Desktop



Docker Networking

- Think of networking in terms of communication boundaries/isolation
 - If two containers are on the same network, they can talk to each other
- Docker runs its own DNS resolver on each network
 - Allows it to resolve IP addresses of other containers using “aliases”



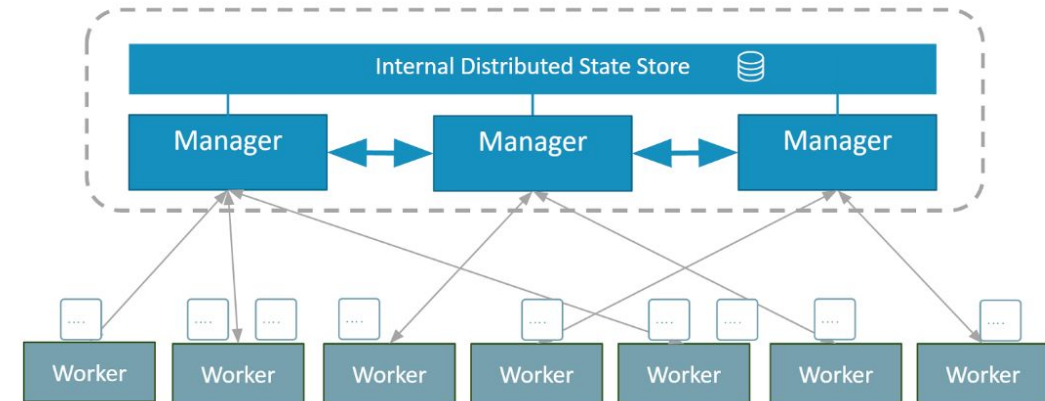
Quick compose demo!

Container Orchestration

- Orchestration provides the ability to manage the running of container workloads, often over a fleet of machines
- You define the **expected state** (the desired state)
- The system then tries to make **actual state** reflect expected state

Actors in Orchestration

- Every orchestrator has the concept of two types of nodes
- Managers
 - Serve as the brains of the cluster
 - Maintain state and schedule work
 - Sometimes called masters
- Worker nodes
 - Perform the actual work, as instructed by a manager
 - Sometimes called agents or nodes



Various Orchestrators

- Docker Swarm
 - Shipped with the Docker engine
 - Very user friendly and easy to get up and running
 - Satisfies most needs, though not all; built to be extensible, but takes some work
- Kubernetes
 - Spun out of work done within Google and contributed to CNCF
 - Think of it more as a toolkit - so not as easy to get up and running
 - Very configurable and extensible
- Amazon ECS
 - Made by Amazon Web Services and provided for free
 - Provides deep integration with AWS resources (IAM, ALBs, Auto-scaling, etc.)

Quick Swarm Demo!

Recap

- Containers/images are here to standardize application packaging
 - No longer require host configuration
 - Docker Compose builds on the abstraction to make multi-service apps easier
 - Container orchestration builds on this idea
- Be mindful of how you build your images and what you include
- Volumes allow data to be persisted longer than the container
- Networking serves provides communication paths/isolation

WARNING!

Containers are NOT a silver bullet that will fix your company culture



Thank you! Rate the session!

Keep in touch!

@mikesir87; mikesir87@vt.edu